

# Data Wrangling R

Art Tay

## Loading Packages

```
# Option 1: Individual Packages (not recommended).
library(dplyr)

# Option 2: tidyverse (okay).
library(tidyverse)

# Option 3: tidymodels (preferred).
library(tidymodels)
tidymodels_prefer() # Can be used to avoid conflicts with other packages.
```

- [Tidyverse Documentation](#)
- [Tidymodels Documentation](#)
- [Data Science in R \(Free Textbook\)](#)

## Pipes

- Pipes are a cleaner way to perform data operations.
- New in R >=4.1 “|>” is a native pipe operator.
- %>% comes from tidyverse and can also be used.
- **Syntax:** object pipe function(args) <-> function(object, args)
- **Chaining:** Data |> function1(args) |> function2(args) = function2(function1(data, args1), args2)

## Examples:

```
# Simple example.
a <- 5
b <- a |> sum(2)
b
```

```
[1] 7
```

```
c <- sum(a, 2)
c
```

```
[1] 7
```

```
# Pipes work with all named functions.
t_test <- rnorm(n = 10, mean = 1, sd = 1) |> t.test()
t_test
```

One Sample t-test

```
data:  rnorm(n = 10, mean = 1, sd = 1)
t = 3.5964, df = 9, p-value = 0.005781
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 0.4220256 1.8530964
sample estimates:
mean of x
 1.137561
```

```
# You can use anonymous functions.
d <- 5 |> {\(x) x * 7}()
d
```

```
[1] 35
```

## Practice:

- Generate 10 samples from a normal distribution, add 1, then test if the mean is different from 2.

```
t_test_practice <- (rnorm(n = 10) + 1) |> t.test(mu = 2)

t_test_practice_2 <- rnorm(n = 10) |> {\(x) x + 1}() |> t.test(mu = 2)
```

## Quick Look at the Data

```
data <- read.csv("./Cholesterol_R.csv")
data <- data |> rename("ID" = contains("ID"))
```

```
head(data)
```

	ID	Before	After4weeks	After8weeks	Margarine
1	1	6.42	5.83	5.75	B
2	2	6.76	6.20	6.13	A
3	3	6.56	5.83	5.71	B
4	4	4.80	4.27	4.15	A
5	5	8.43	7.71	7.67	B
6	6	7.49	7.12	7.05	A

```
str(data)
```

```
'data.frame':  18 obs. of  5 variables:
 $ ID          : int  1 2 3 4 5 6 7 8 9 10 ...
 $ Before      : num  6.42 6.76 6.56 4.8 8.43 7.49 8.05 5.05 5.77 3.91 ...
 $ After4weeks: num  5.83 6.2 5.83 4.27 7.71 7.12 7.25 4.63 5.31 3.7 ...
 $ After8weeks: num  5.75 6.13 5.71 4.15 7.67 7.05 7.1 4.67 5.33 3.66 ...
 $ Margarine   : chr  "B" "A" "B" "A" ...
```

## Subsetting

### Examples:

```
# Selecting columns by name.
measurements <- data |> select(Before, After4weeks, After8weeks)
head(measurements)
```

	Before	After4weeks	After8weeks
1	6.42	5.83	5.75
2	6.76	6.20	6.13
3	6.56	5.83	5.71
4	4.80	4.27	4.15
5	8.43	7.71	7.67
6	7.49	7.12	7.05

```
# Selecting columns based on a pattern.
after <- data |> select(starts_with("After"))
head(after)
```

	After4weeks	After8weeks
1	5.83	5.75
2	6.20	6.13
3	5.83	5.71
4	4.27	4.15
5	7.71	7.67
6	7.12	7.05

```
# Lots of options.
tidyselect::starts_with()
tidyselect::ends_with()
tidyselect::contains()
tidyselect::matches()
tidyselect::num_range()
tidyselect::everything()
tidyselect::one_of()
tidyselect::all_of()
tidyselect::any_of()
```

```
# Slicing rows.
first3 <- data |> slice(1:3)
first3
```

	ID	Before	After4weeks	After8weeks	Margarine
1	1	6.42	5.83	5.75	B
2	2	6.76	6.20	6.13	A
3	3	6.56	5.83	5.71	B

```
# Filter rows based on a condition.
A_above_6 <- data |> filter(After8weeks >= 6 & Margarine == "A")
head(A_above_6)
```

	ID	Before	After4weeks	After8weeks	Margarine
--	----	--------	-------------	-------------	-----------

1	2	6.76	6.20	6.13	A
2	6	7.49	7.12	7.05	A
3	14	7.67	7.11	6.96	A
4	15	7.34	6.84	6.82	A

## Practice

- Find the patient IDs with baseline (before) measurements below 5.

```
data |> filter(Before < 5) |> select(ID)
```

```
ID
1  4
2 10
```

## Sorting

### Examples

```
# Sort the data based on a single column.
data_sort_baseline <- data |> arrange(desc(Before))
head(data_sort_baseline)
```

	ID	Before	After4weeks	After8weeks	Margarine
1	5	8.43	7.71	7.67	B
2	7	8.05	7.25	7.10	B
3	14	7.67	7.11	6.96	A
4	6	7.49	7.12	7.05	A
5	15	7.34	6.84	6.82	A
6	16	6.85	6.40	6.29	B

```
# Sort thr data based on multiple columns.
data_sort_many <- data |> arrange(Before, Margarine)
head(data_sort_many)
```

	ID	Before	After4weeks	After8weeks	Margarine
1	10	3.91	3.70	3.66	A
2	4	4.80	4.27	4.15	A
3	8	5.05	4.63	4.67	A
4	17	5.13	4.52	4.45	A
5	18	5.73	5.13	5.17	B
6	9	5.77	5.31	5.33	B

## Practice

- Sort Margarine in alphabetical order, then the IDs in ascending order.

```
data_sort_practice <- data |> arrange(Margarine, ID)
data_sort_practice
```

	ID	Before	After4weeks	After8weeks	Margarine
1	2	6.76	6.20	6.13	A
2	4	4.80	4.27	4.15	A

3	6	7.49	7.12	7.05	A
4	8	5.05	4.63	4.67	A
5	10	3.91	3.70	3.66	A
6	13	6.17	5.56	5.51	A
7	14	7.67	7.11	6.96	A
8	15	7.34	6.84	6.82	A
9	17	5.13	4.52	4.45	A
10	1	6.42	5.83	5.75	B
11	3	6.56	5.83	5.71	B
12	5	8.43	7.71	7.67	B
13	7	8.05	7.25	7.10	B
14	9	5.77	5.31	5.33	B
15	11	6.77	6.15	5.96	B
16	12	6.44	5.59	5.64	B
17	16	6.85	6.40	6.29	B
18	18	5.73	5.13	5.17	B

## Summarizing

### Example

```
# Creating a table of summary statistics.
summary_table <- data |> select(-ID) |> group_by(Margarine) |>
  summarize(
    across(where(is.numeric),
      list(mean = mean, sd = sd)
    )
  )

summary_table
```

```
# A tibble: 2 x 7
  Margarine Before_mean Before_sd After4weeks_mean After4weeks_sd
  <chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 A          6.04      1.36      5.55      1.32
2 B          6.78      0.919     6.13      0.864
# i 2 more variables: After8weeks_mean <dbl>, After8weeks_sd <dbl>
```

### Practice

- Find the median for each column for patients with even ID numbers.

```
summary_table_practice <- data |> filter(ID %% 2 == 0) |> select(-ID) |>
  summarize(
    across(where(is.numeric),
      list(median = median)
    )
  )
```

# Feature Engineering

## Examples

```
# Create simple features.
data <- data |> mutate(
  diff_4wk = After4weeks - Before,
  diff_8wk = After8weeks - Before
)

head(data)
```

	ID	Before	After4weeks	After8weeks	Margarine	diff_4wk	diff_8wk
1	1	6.42	5.83	5.75	B	-0.59	-0.67
2	2	6.76	6.20	6.13	A	-0.56	-0.63
3	3	6.56	5.83	5.71	B	-0.73	-0.85
4	4	4.80	4.27	4.15	A	-0.53	-0.65
5	5	8.43	7.71	7.67	B	-0.72	-0.76
6	6	7.49	7.12	7.05	A	-0.37	-0.44

```
# Function based example: normalization.
norm_func <- function(x){
  x <- x - mean(x)
  return(x / sd(x))
}

normalized <- data |> select(Before) |>
  mutate(
    across(where(is.numeric), list(norm = norm_func, log = log))
  )

head(normalized)
```

	Before	Before_norm	Before_log
1	6.42	0.0102614	1.859418
2	6.76	0.2957149	1.911023
3	6.56	0.1278011	1.880991
4	4.80	-1.3498404	1.568616
5	8.43	1.6977952	2.131797
6	7.49	0.9086003	2.013569

## Practice

- Find the percentage change between baseline and 8 weeks.

```
data_percent_change <- data |> mutate(percent_change = After8weeks / Before - 1)
```

# Pivoting


## Examples

```
# Pivot to long format. (Good for plotting).
data_long <- data |>
  pivot_longer(cols = c(Before, starts_with("After")),
               names_to = "Time", values_to = "Measurements")
head(data_long)
```

# A tibble: 6 x 6

	ID	Margarine	diff_4wk	diff_8wk	Time	Measurements
	<int>	<chr>	<dbl>	<dbl>	<chr>	<dbl>
1	1	B	-0.59	-0.67	Before	6.42
2	1	B	-0.59	-0.67	After4weeks	5.83
3	1	B	-0.59	-0.67	After8weeks	5.75
4	2	A	-0.560	-0.63	Before	6.76
5	2	A	-0.560	-0.63	After4weeks	6.2
6	2	A	-0.560	-0.63	After8weeks	6.13

```
# Why?
data_long |> ggplot(aes(x = ID, y = Measurements, color = Time)) +
  geom_point() + geom_line()
```



```
# Pivot to wide format.
data_wide <- data_long |> pivot_wider(names_from = Time,
                                       values_from = Measurements)
head(data_wide)
```

# A tibble: 6 x 7

	ID	Margarine	diff_4wk	diff_8wk	Before	After4weeks	After8weeks
	<int>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1	B	-0.59	-0.67	6.42	5.83	5.75
2	2	A	-0.560	-0.63	6.76	6.2	6.13
3	3	B	-0.730	-0.85	6.56	5.83	5.71
4	4	A	-0.53	-0.650	4.8	4.27	4.15
5	5	B	-0.72	-0.76	8.43	7.71	7.67
6	6	A	-0.37	-0.440	7.49	7.12	7.05

```
# Why?
t.test(data_wide$Before, data_wide$After8weeks)
```

Welch Two Sample t-test

```
data: data_wide$Before and data_wide$After8weeks
t = 1.6443, df = 33.796, p-value = 0.1094
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-0.1485269 1.4063046
sample estimates:
mean of x mean of y
6.407778 5.778889
```

## Practice

- Create a bar plot with time groups on the x axis and the average 8 week difference on the y axis.

```
plot_practice <- data |>
  pivot_longer(cols = c(Before, starts_with("After")),
    names_to = "Time", values_to = "Measurements") |>
  group_by(Time) |>
  summarise(avg_8_diff = mean(Measurements)) |>
  ggplot(aes(x = Time, y = avg_8_diff)) +
  geom_bar(stat = "identity")
```

## Joins

### Examples

```
# Include additional information.
info_A <- data |> select(ID, Margarine) |> filter(Margarine == "A") |>
  mutate(After12weeks = rnorm(length(ID), mean = 3, sd = 1))
head(info_A)
```

	ID	Margarine	After12weeks
1	2	A	2.026017
2	4	A	2.928692
3	6	A	1.906760
4	8	A	2.139485
5	10	A	2.237223
6	13	A	2.745113



```
# Left join - maintains all the information on the left table.
data_left <- data |> left_join(info_A, join_by(ID, Margarine))
head(data_left)
```

	ID	Before	After4weeks	After8weeks	Margarine	diff_4wk	diff_8wk	After12weeks
1	1	6.42	5.83	5.75	B	-0.59	-0.67	NA
2	2	6.76	6.20	6.13	A	-0.56	-0.63	2.026017
3	3	6.56	5.83	5.71	B	-0.73	-0.85	NA
4	4	4.80	4.27	4.15	A	-0.53	-0.65	2.928692
5	5	8.43	7.71	7.67	B	-0.72	-0.76	NA
6	6	7.49	7.12	7.05	A	-0.37	-0.44	1.906760

```
# Right join - maintains all the information on the right table.
data_right <- data |> right_join(info_A, join_by(ID, Margarine))
head(data_right)
```

	ID	Before	After4weeks	After8weeks	Margarine	diff_4wk	diff_8wk	After12weeks
1	2	6.76	6.20	6.13	A	-0.56	-0.63	2.026017
2	4	4.80	4.27	4.15	A	-0.53	-0.65	2.928692
3	6	7.49	7.12	7.05	A	-0.37	-0.44	1.906760
4	8	5.05	4.63	4.67	A	-0.42	-0.38	2.139485
5	10	3.91	3.70	3.66	A	-0.21	-0.25	2.237223
6	13	6.17	5.56	5.51	A	-0.61	-0.66	2.745113

```
# Inner v. Full Joins
full_ex_1 <- data.frame(ID = 1:10, value_1 = rnorm(10))
full_ex_2 <- data.frame(ID = 6:15, value_2 = rnorm(10))

full_ex_1 |> left_join(full_ex_2, join_by(ID))
```

	ID	value_1	value_2
1	1	-0.33878578	NA
2	2	-0.47538017	NA
3	3	-0.28835249	NA
4	4	2.33931073	NA
5	5	-1.36536669	NA
6	6	0.49142987	-0.05060684
7	7	0.47420285	0.36681646
8	8	-0.57437554	0.54348327
9	9	0.73461252	0.08616754
10	10	0.03344151	0.65897144

```
full_ex_1 |> inner_join(full_ex_2, join_by(ID))
```

	ID	value_1	value_2
1	6	0.49142987	-0.05060684
2	7	0.47420285	0.36681646
3	8	-0.57437554	0.54348327
4	9	0.73461252	0.08616754
5	10	0.03344151	0.65897144

```
full_ex_1 |> full_join(full_ex_2, join_by(ID))
```

	ID	value_1	value_2
1	1	-0.33878578	NA
2	2	-0.47538017	NA
3	3	-0.28835249	NA
4	4	2.33931073	NA
5	5	-1.36536669	NA
6	6	0.49142987	-0.05060684
7	7	0.47420285	0.36681646
8	8	-0.57437554	0.54348327
9	9	0.73461252	0.08616754
10	10	0.03344151	0.65897144
11	11	NA	0.01532030
12	12	NA	0.46015989
13	13	NA	-0.24378432
14	14	NA	-0.89058840
15	15	NA	1.69796996

```
# Additional patients to include.
data_2 <- data |> select(-c(ID, starts_with("diff"))) |>
  mutate(Margarine = ifelse(Margarine == "A", "C", "D")) |>
  mutate(across(where(is.numeric),
    function(x){x + rnorm(length(x))}
  )
) |>
  mutate(ID = 19:36)

data_combined <- data |> bind_rows(data_2)
head(data_combined)
```

	ID	Before	After4weeks	After8weeks	Margarine	diff_4wk	diff_8wk
1	1	6.42	5.83	5.75	B	-0.59	-0.67
2	2	6.76	6.20	6.13	A	-0.56	-0.63
3	3	6.56	5.83	5.71	B	-0.73	-0.85
4	4	4.80	4.27	4.15	A	-0.53	-0.65
5	5	8.43	7.71	7.67	B	-0.72	-0.76
6	6	7.49	7.12	7.05	A	-0.37	-0.44

## Practice

- Include the following information retaining all patients present in the original dataset measurement. Complete the dataset by replacing any missing (NA) values with the mean for the patients corresponding group (Margarine).

```
data_to_include <- data.frame(
  ID = 6:20,
  Weight = rnorm(n = 15, 100, 10)
)

data_practice <- data |> left_join(data_to_include) |>
  group_by(Margarine) |>
  mutate(Weight = ifelse(is.na(Weight), mean(Weight, na.rm = T), Weight))
```

## Extensions for Modeling

### Resampling

```
# Create a 80-20 train/test split.
set.seed(123)
split <- initial_split(data_combined, prop = 0.8) # contains indices.

data_train <- training(split)
data_test <- testing(split)
```

```
# Bootstrap t-tests.
boots <- data_train |> select(ID, Before) |> bootstraps(times = 10)
head(as.data.frame(boots$splits[[1]]))
```

```
ID   Before
1  5 8.430000
2  5 8.430000
3 20 6.761933
4  4 4.800000
5 24 6.308628
6 12 6.440000
```

```
p_values <- sapply(boots$splits,
  FUN = function(x) {
    p_value <- x |> as.data.frame() |> t.test()
    p_value$p.value
  }
)
p_values
```

```
[1] 1.990771e-14 2.714639e-14 9.917317e-16 6.688158e-13 9.027094e-14
[6] 1.217726e-14 6.462299e-14 3.062132e-14 3.848936e-15 2.807957e-13
```

### Recipes

- Recipes are a way to bundle data pre-processing (possible trained).
- Easier access to more complex transformations (ex: PCA, Splines, Imputation).
- Avoids data leakage, and is online for incoming data.
- Works well with the resampling and modeling framework.

```
# Recipe work based on variable roles. This makes it easier to preform
# grouped transformations.
recipe_ex <- data_train |>
  recipe(diff_8wk ~ .) |>
  # Don't use ID for modeling, but keep it in the dataset.
  update_role(ID, new_role = "ID")

recipe_ex
```

```
# Most dplyr function have an analog in recipes.
recipe_ex <- recipe_ex |>
  step_select(c(diff_8wk, Margarine, Before)) |>
  step_mutate(Margarine = as.factor(Margarine))
```

```
recipe_ex
```

```
# Let's take a look at what our cleaned dataset would look like here.
data_train_cleaned <- bake(prepare(recipe_ex), new_data = NULL)
head(data_train_cleaned)
```

```
# A tibble: 6 x 3
  diff_8wk Margarine Before
    <dbl>   <fct>     <dbl>
1    NA     C         7.45
2  -0.520 A         7.34
3  -0.71  A         7.67
4  -0.85  B         6.56
5  -0.25  A         3.91
6  -0.560 B         5.73
```

```
# Why do we need prep? For learned transformations.
recipe_ex <- recipe_ex |>
  # Imputation - handles missing data.
  step_impute_knn(diff_8wk, neighbors = 1)
```

```
recipe_ex |> prep()
```

```
# Grouped transformations.
recipe_ex <- recipe_ex |>
  # Normalize - there is actually a step_normalize.
  step_mutate_at(all_numeric_predictors(), fn = norm_func) |>
  # Create dummy variable.
  step_dummy(all_nominal_predictors()) |>
  prep()
```

```
recipe_ex
```

```
data_train_cleaned <- bake(recipe_ex, new_data = NULL)
data_train_cleaned
```

```
# A tibble: 28 x 5
  diff_8wk Before Margarine_B Margarine_C Margarine_D
    <dbl>   <dbl>     <dbl>     <dbl>     <dbl>
1  -0.520  0.768         0         1         0
2  -0.520  0.693         0         0         0
3  -0.71   0.922         0         0         0
4  -0.85   0.151         1         0         0
5  -0.25  -1.69         0         0         0
6  -0.560 -0.425         1         0         0
7  -0.25  -1.39         0         1         0
8  -0.810  0.297         1         0         0
```

```

  9   -0.76   1.45           1           0           0
10   -0.810  0.291          0           1           0
# i 18 more rows

data_test_cleaned <- bake(recipe_ex, new_data = data_test)
data_test_cleaned

# A tibble: 8 x 5
  diff_8wk Before Margarine_B Margarine_C Margarine_D
  <dbl>   <dbl>         <dbl>         <dbl>         <dbl>
1  -0.63   0.260           0           0           0
2  -0.440  1.10            0           0           0
3  -0.66  -0.420           0           0           0
4  -0.560  0.364           1           0           0
5  -0.68  -1.62           0           0           0
6  -0.71   1.29           0           0           1
7  -0.560 -1.05           0           1           0
8  -0.85   0.0714        0           0           1

```